

With TF 1.0!

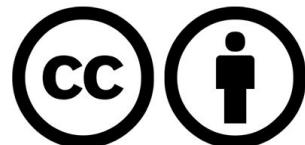


# Lab 6

## Softmax Classifier

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



# Call for comments

Please feel free to add comments directly on these slides

Other slides: <https://goo.gl/jPtWNt>



With TF 1.0!

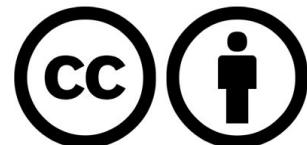


# Lab 6-1

## Softmax Classifier

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>

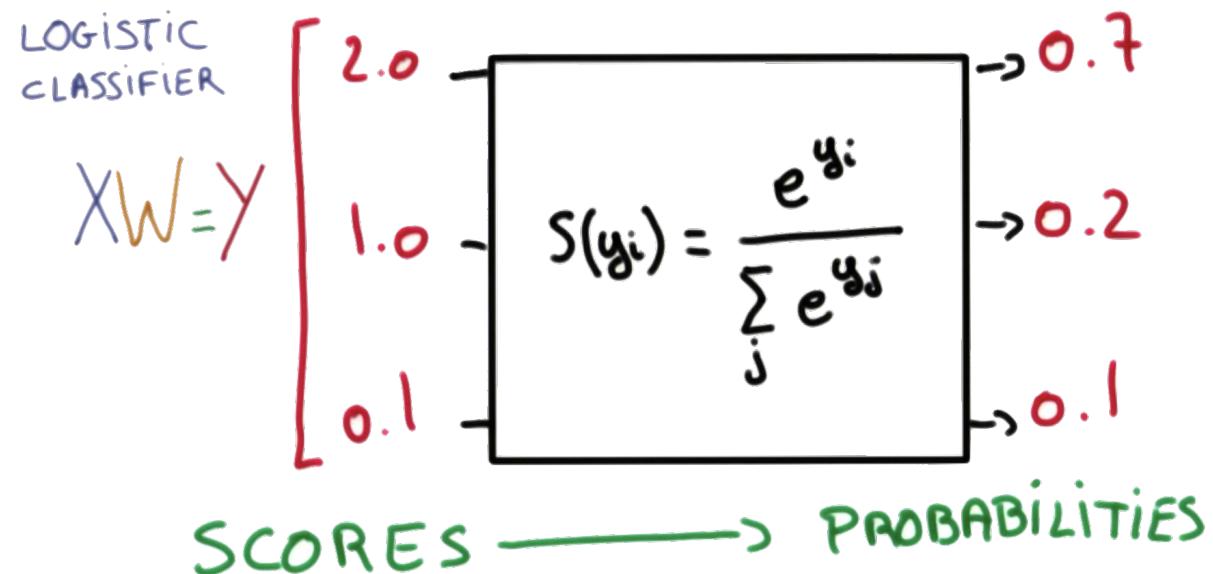
Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/>

 <b>zeran4</b> 1 commit / 5 ++ / 4 --	#11	 <b>jennykang</b> 19 commits / 940 ++ / 253 --	#2
 <b>GzuPark</b> 14 commits / 41 ++ / 31 --	#3	 <b>kkweon</b> 5 commits / 372 ++ / 296 --	#4
 <b>BlueMelon715</b> 4 commits / 45 ++ / 34 --	#5	 <b>jin-chong</b> 2 commits / 4 ++ / 4 --	#6
 <b>FuZer</b> 2 commits / 37 ++ / 30 --	#7	 <b>cynthia</b> 1 commit / 28 ++ / 28 --	#8
 <b>keon</b> 1 commit / 3 ++ / 3 --	#9	 <b>allieus</b> 1 commit / 55 ++ / 59 --	#10

# Softmax function



hypothesis = tf.nn.softmax(tf.matmul(X,W)+b)



# Cost function: cross entropy

$$\text{LOSS} \downarrow$$
$$\mathcal{L} = \frac{1}{N} \sum_i D(s(w_i x_i + b), L_i)$$

TRAINING SET

STEP

$$-\alpha \Delta \mathcal{L}(w_1, w_2)$$

DERIVATIVE

# Cross entropy cost/loss

```
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```

```
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)
```

```

x_data = [[1, 2, 1, 1], [2, 1, 3, 2], [3, 1, 3, 4], [4, 1, 5, 5], [1, 7, 5, 5],
           [1, 2, 5, 6], [1, 6, 6, 6], [1, 7, 7, 7]]
y_data = [[0, 0, 1], [0, 0, 1], [0, 0, 1], [0, 1, 0], [0, 1, 0], [0, 1, 0], [1, 0, 0], [1, 0, 0]]

X = tf.placeholder("float", [None, 4])
Y = tf.placeholder("float", [None, 3])
nb_classes = 3

W = tf.Variable(tf.random_normal([4, nb_classes]), name='weight')
b = tf.Variable(tf.random_normal([nb_classes]), name='bias')

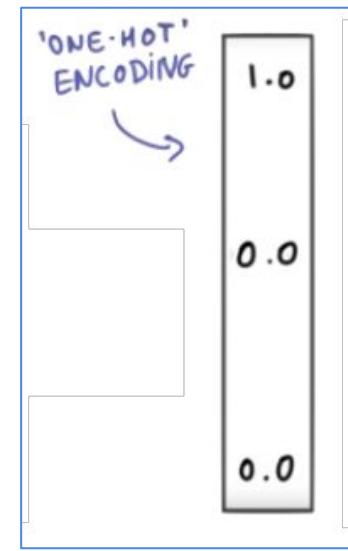
# tf.nn.softmax computes softmax activations
# softmax = exp(logits) / reduce_sum(exp(logits), dim)
hypothesis = tf.nn.softmax(tf.matmul(X, W) + b)

# Cross entropy cost/loss
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)

# Launch graph
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())

    for step in range(2001):
        sess.run(optimizer, feed_dict={X: x_data, Y: y_data})
        if step % 200 == 0:
            print(step, sess.run(cost, feed_dict={X: x_data, Y: y_data}))

```



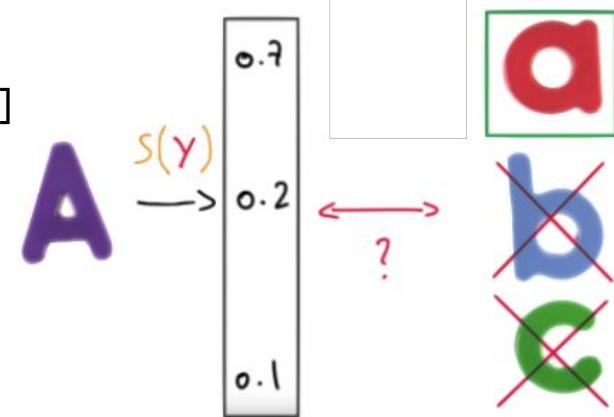
# Test & one-hot encoding

```
hypothesis = tf.nn.softmax(tf.matmul(X,W)+b)
```

# Testing & One-hot encoding

```
a = sess.run(hypothesis, feed_dict={X: [[1, 11, 7, 9]]})  
print(a, sess.run(tf.argmax(a, 1)))
```

```
[[ 1.38904958e-03  9.98601854e-01  9.06129117e-06]] [1]
```



# Test & one-hot encoding

```
hypothesis = tf.nn.softmax(tf.matmul(X,W)+b)
```

```
all = sess.run(hypothesis, feed_dict={X: [[1, 11, 7, 9],  
                                         [1, 3, 4, 3],  
                                         [1, 1, 0, 1]]})  
print(all, sess.run(tf.argmax(all, 1)))
```

```
[[ 1.38904958e-03  9.98601854e-01  9.06129117e-06]  
[ 9.31192040e-01  6.29020557e-02  5.90589503e-03]  
[ 1.27327668e-08  3.34112905e-04  9.99665856e-01]]
```

[1 0 2]

With TF 1.0!



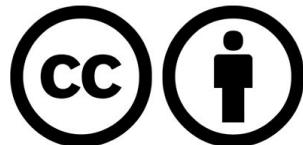
# Lab 6-2

## Fancy Softmax Classifier

*cross\_entropy, one\_hot, reshape*

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/>

 <b>zeran4</b> 1 commit / 5 ++ / 4 --	#11	 <b>jennykang</b> 19 commits / 940 ++ / 253 --	#2
 <b>GzuPark</b> 14 commits / 41 ++ / 31 --	#3	 <b>kkweon</b> 5 commits / 372 ++ / 296 --	#4
 <b>BlueMelon715</b> 4 commits / 45 ++ / 34 --	#5	 <b>jin-chong</b> 2 commits / 4 ++ / 4 --	#6
 <b>FuZer</b> 2 commits / 37 ++ / 30 --	#7	 <b>cynthia</b> 1 commit / 28 ++ / 28 --	#8
 <b>keon</b> 1 commit / 3 ++ / 3 --	#9	 <b>allieus</b> 1 commit / 55 ++ / 59 --	#10

# softmax\_cross\_entropy\_with\_logits

```
logits = tf.matmul(X, W) + b  
hypothesis = tf.nn.softmax(logits)
```

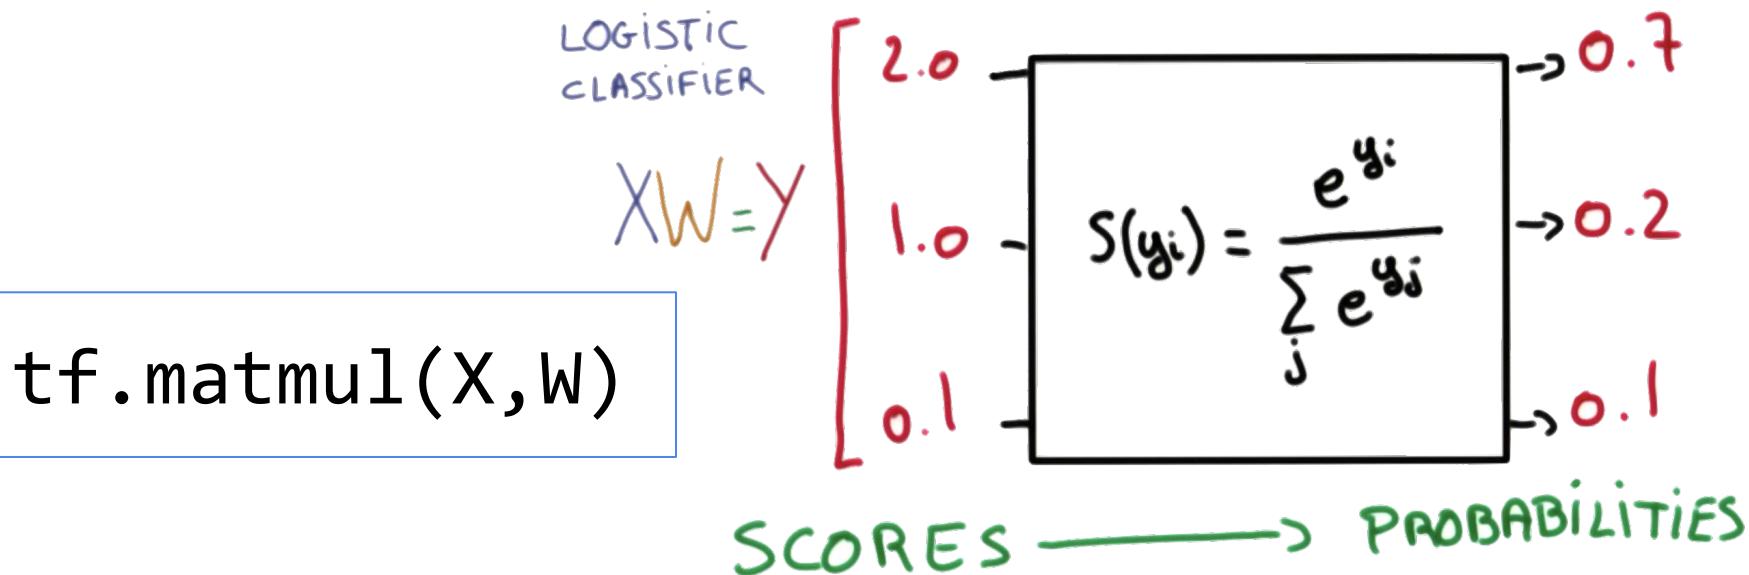
1

```
# Cross entropy cost/loss  
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```

2

```
"Cross entropy cost/loss  
cost_i = tf.nn.softmax_cross_entropy_with_logits(logits=logits,  
                                                labels=Y_one_hot)  
cost = tf.reduce_mean(cost_i)
```

hypothesis = tf.nn.softmax(tf.matmul(X,W))



# softmax\_cross\_entropy\_with\_logits

```
logits = tf.matmul(X, W) + b  
hypothesis = tf.nn.softmax(logits)
```

1

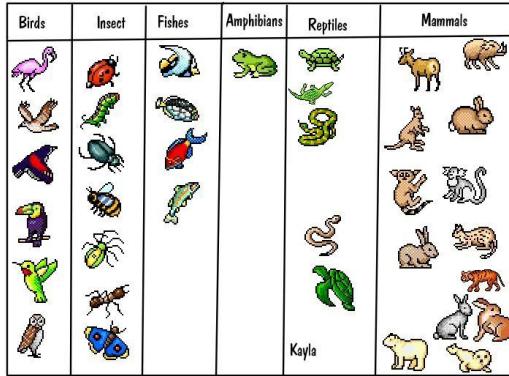
```
# Cross entropy cost/loss  
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```

2

```
"Cross entropy cost/loss  
cost_i = tf.nn.softmax_cross_entropy_with_logits(logits=logits,  
                                                labels=Y_one_hot)  
cost = tf.reduce_mean(cost_i)
```

# Animal classification

## with softmax\_cross\_entropy\_with\_logits



1	0	0	1	0	0	0	1	1	1	0	0	4	1	0	1	0
0	0	1	0	0	1	1	1	1	0	0	1	0	1	0	0	3
1	0	0	1	0	0	1	1	1	1	0	0	4	0	0	1	0
1	0	0	1	0	0	1	1	1	1	0	0	4	1	0	1	0
1	0	0	1	0	0	0	1	1	1	0	0	0	4	1	0	1
1	0	0	1	0	0	0	1	1	1	0	0	0	4	1	1	0
0	0	1	0	0	1	0	1	1	1	0	0	1	0	1	1	0
0	0	1	0	0	1	1	1	1	1	0	0	1	0	1	0	3
0	0	1	0	0	1	0	1	1	1	0	0	1	0	1	0	3
1	0	0	1	0	0	0	1	1	1	1	0	0	0	4	0	1
1	0	0	1	0	0	0	1	1	1	1	0	0	0	4	1	0
0	1	1	0	0	1	0	0	1	1	0	0	0	2	1	1	0
0	0	1	0	0	0	1	1	1	1	0	0	0	1	0	0	3
0	0	1	0	0	0	0	1	0	0	0	0	0	0	0	0	6
0	0	1	0	0	0	1	1	0	0	0	0	0	4	0	0	6
0	0	1	0	0	0	1	1	0	0	0	0	0	6	0	0	6
0	1	1	0	1	0	1	0	1	0	0	0	0	2	1	0	1
1	0	0	1	0	0	0	1	1	1	1	0	0	4	1	0	1

# Predicting animal type based on various features

```
xy = np.loadtxt('data-04-zoo.csv', delimiter=',', dtype=np.float32)
x_data = xy[:, 0:-1]
y_data = xy[:, [-1]]
```

# tf.one\_hot and reshape

1	0	0	1	0	0	0	1	1	1	0	0	4	1	0	1	0
0	0	1	0	0	1	1	1	1	0	0	1	0	1	0	0	3
1	0	0	1	0	0	1	1	1	1	0	0	4	0	0	1	0
1	0	0	1	0	0	1	1	1	1	0	0	4	1	0	1	0
1	0	0	1	0	0	1	1	1	1	0	0	4	1	0	1	0
1	0	0	1	0	0	1	1	1	1	0	0	4	1	1	1	0
0	0	1	0	0	1	0	1	1	0	0	1	0	1	1	0	3
0	0	1	0	0	1	1	1	1	0	0	1	0	1	0	0	3
1	0	0	1	0	0	1	1	1	1	0	0	4	0	1	0	0
1	0	0	1	0	0	1	1	1	1	0	0	4	1	0	1	0
0	1	1	0	1	0	0	0	1	1	0	0	2	1	1	0	1
0	0	1	0	0	1	1	1	1	0	0	1	0	1	0	0	3
0	0	1	0	0	0	1	0	0	0	0	0	0	0	0	0	6
0	0	1	0	0	1	1	0	0	0	0	0	4	0	0	0	6
0	0	1	0	0	1	1	0	0	0	0	0	6	0	0	0	6
0	1	1	0	1	0	1	0	1	1	0	0	2	1	0	0	1
1	0	0	1	0	0	0	1	1	1	0	0	4	1	0	1	0

```
Y = tf.placeholder(tf.int32, [None, 1]) # 0 ~ 6, shape=(?, 1)
Y_one_hot = tf.one_hot(Y, nb_classes) # one hot shape=(?, 1, 7)
Y_one_hot = tf.reshape(Y_one_hot, [-1, nb_classes]) # shape=(?, 7)
```

If the input indices is rank N, the output will have rank N+1. The new axis is created at dimension axis (default: the new axis is appended at the end).

[https://www.tensorflow.org/api\\_docs/python/tf/one\\_hot](https://www.tensorflow.org/api_docs/python/tf/one_hot)

```

# Predicting animal type based on various features
xy = np.loadtxt('data-04-zoo.csv', delimiter=',', dtype=np.float32)
x_data = xy[:, 0:-1]
y_data = xy[:, [-1]]

nb_classes = 7 # 0 ~ 6

X = tf.placeholder(tf.float32, [None, 16])
Y = tf.placeholder(tf.int32, [None, 1]) # 0 ~ 6

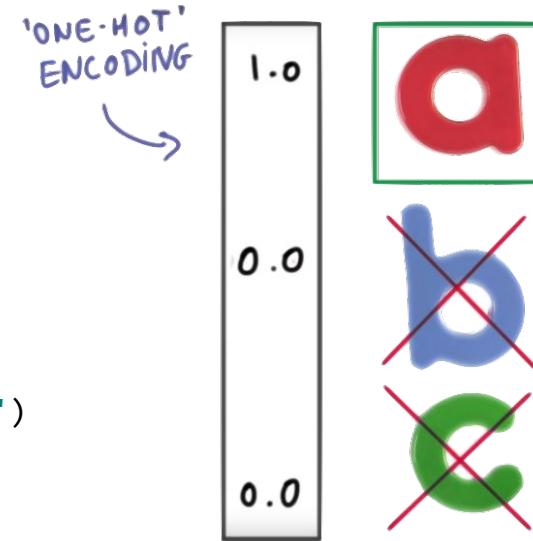
Y_one_hot = tf.one_hot(Y, nb_classes) # one hot
Y_one_hot = tf.reshape(Y_one_hot, [-1, nb_classes])

W = tf.Variable(tf.random_normal([16, nb_classes]), name='weight')
b = tf.Variable(tf.random_normal([nb_classes]), name='bias')

# tf.nn.softmax computes softmax activations
# softmax = exp(Logits) / reduce_sum(exp(Logits), dim)
logits = tf.matmul(X, W) + b
hypothesis = tf.nn.softmax(logits)

# Cross entropy cost/loss
cost_i = tf.nn.softmax_cross_entropy_with_logits(logits=logits,
                                                labels=Y_one_hot)
cost = tf.reduce_mean(cost_i)
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)

```



```

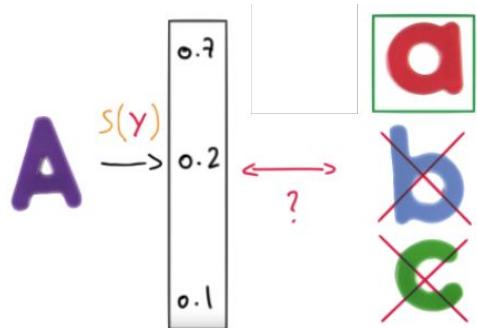
cost = tf.reduce_mean(cost_i)
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)

prediction = tf.argmax(hypothesis, 1)
correct_prediction = tf.equal(prediction, tf.argmax(Y_one_hot, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
# Launch graph
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())

    for step in range(2000):
        sess.run(optimizer, feed_dict={X: x_data, Y: y_data})
        if step % 100 == 0:
            loss, acc = sess.run([cost, accuracy], feed_dict={
                X: x_data, Y: y_data})
            print("Step: {:5}\tLoss: {:.3f}\tAcc: {:.2%}".format(
                step, loss, acc))

    # Let's see if we can predict
    pred = sess.run(prediction, feed_dict={X: x_data})
    # y_data: (N,1) = flatten => (N, ) matches pred.shape
    for p, y in zip(pred, y_data.flatten()):
        print("[{}] Prediction: {} True Y: {}".format(p == int(y), p, int(y)))

```



```

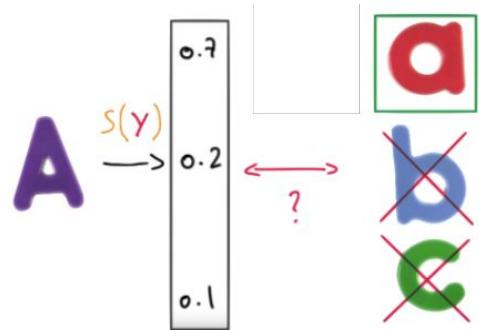
cost = tf.reduce_mean(cost_i)
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)

prediction = tf.argmax(hypothesis, 1)
correct_prediction = tf.equal(prediction, tf.argmax(Y_one_hot, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
# Launch graph
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())

    for step in range(2000):
        sess.run(optimizer, feed_dict={X: x_data, Y: y_data})
        if step % 100 == 0:
            loss, acc = sess.run([cost, accuracy], feed_dict={
                X: x_data, Y: y_data})
            print("Step: {:5}\tLoss: {:.3f}\tAcc: {:.2%}".format(
                step, loss, acc))

    # Let's see if we can predict
    pred = sess.run(prediction, feed_dict={X: x_data})
    # y_data: (N,1) = flatten => (N, ) matches pred.shape
    for p, y in zip(pred, y_data.flatten()):
        print("[{}]\tPrediction: {} True Y: {}".format(
            p, int(y), int(y)))

```



Step: 1100 Loss: 0.101 Acc: 99.01%  
Step: 1200 Loss: 0.092 Acc: 100.00%  
Step: 1300 Loss: 0.084 Acc: 100.00%

...

[True] Prediction: 0 True Y: 0  
[True] Prediction: 0 True Y: 0  
[True] Prediction: 3 True Y: 3  
[True] Prediction: 0 True Y: 0  
[True] Prediction: 3 True Y: 3  
[True] Prediction: 3 True Y: 3  
[True] Prediction: 0 True Y: 0

# Lab 7

## Learning rate, Evaluation

Sung Kim <hunkim+ml@gmail.com>