

With TF 1.0!




# Lab 7-I

## Learning rate, Evaluation

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



# Call for comments

Please feel free to add comments directly on these slides

Other slides: <https://goo.gl/jPtWNt>



With TF 1.0!

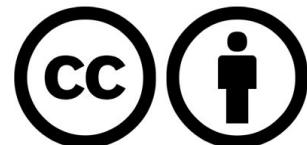


# Lab 7-I

## Learning rate, Evaluation

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/>

 <b>zeran4</b> 1 commit / 5 ++ / 4 --	#11	 <b>jennykang</b> 19 commits / 940 ++ / 253 --	#2
 <b>GzuPark</b> 14 commits / 41 ++ / 31 --	#3	 <b>kkweon</b> 5 commits / 372 ++ / 296 --	#4
 <b>BlueMelon715</b> 4 commits / 45 ++ / 34 --	#5	 <b>jin-chong</b> 2 commits / 4 ++ / 4 --	#6
 <b>FuZer</b> 2 commits / 37 ++ / 30 --	#7	 <b>cynthia</b> 1 commit / 28 ++ / 28 --	#8
 <b>keon</b> 1 commit / 3 ++ / 3 --	#9	 <b>allieus</b> 1 commit / 55 ++ / 59 --	#10

# Training and Test datasets



```
x_data = [[1, 2, 1], [1, 3, 2], [1, 3, 4], [1, 5, 5], [1, 7, 5], [1, 2, 5], [1, 6, 6], [1, 7, 7]]  
y_data = [[0, 0, 1], [0, 0, 1], [0, 0, 1], [0, 1, 0], [0, 1, 0], [0, 1, 0], [1, 0, 0], [1, 0, 0]]  
  
# Evaluation our model using this test dataset  
x_test = [[2, 1, 1], [3, 1, 2], [3, 3, 4]]  
y_test = [[0, 0, 1], [0, 0, 1], [0, 0, 1]]
```

```
X = tf.placeholder("float", [None, 3])
Y = tf.placeholder("float", [None, 3])
W = tf.Variable(tf.random_normal([3, 3]))
b = tf.Variable(tf.random_normal([3]))
```

```
hypothesis = tf.nn.softmax(tf.matmul(X, W)+b)
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)
```

```
# Correct prediction Test model
```

```
prediction = tf.argmax(hypothesis, 1)
is_correct = tf.equal(prediction, tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

```
# Launch graph
```

```
with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())
    for step in range(201):
        cost_val, W_val, _ = sess.run([cost, W, optimizer],
                                      feed_dict={X: x_data, Y: y_data})
        print(step, cost_val, W_val)
```

```
# predict
```

```
print("Prediction:", sess.run(prediction, feed_dict={X: x_test}))
```

```
# Calculate the accuracy
```

```
print("Accuracy: ", sess.run(accuracy, feed_dict={X: x_test, Y: y_test}))
```

# Training and Test datasets

```
199 0.672261 [[-1.15377033 0.28146935  
1.13632679]
```

```
[ 0.37484586 0.18958236 0.33544877]  
[-0.35609841 -0.43973011 -1.25604188]]
```


```
200 0.670909 [[-1.15885413 0.28058422  
1.14229572]
```

```
[ 0.37609792 0.19073224 0.33304682]  
[-0.35536593 -0.44033223 -1.2561723 ]]
```


```
Prediction: [2 2 2]
```

```
Accuracy: 1.0
```

# Learning rate: NaN!



Large learning rate: Overshooting.



Small learning rate: Many iterations until convergence and trapping in local minima.

# Big learning rate

```
X = tf.placeholder("float", [None, 3])
Y = tf.placeholder("float", [None, 3])
W = tf.Variable(tf.random_normal([3, 3]))
b = tf.Variable(tf.random_normal([3]))
hypothesis = tf.nn.softmax(tf.matmul(X, W)+b)
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=1.5).minimize(cost)

# Correct prediction Test model
prediction = tf.argmax(hypothesis, 1)
is_correct = tf.equal(prediction, tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))

# Launch graph
with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())
    for step in range(201):
        cost_val, W_val, _ = sess.run([cost, W, optimizer],
                                      feed_dict={X: x_data, Y: y_data})
        print(step, cost_val, W_val)

    # predict
    print("Prediction:", sess.run(prediction, feed_dict={X: x_test}))
    # Calculate the accuracy
    print("Accuracy: ", sess.run(accuracy, feed_dict={X: x_test, Y: y_test}))
```

```
2 27.2798 [[ 0.44451016 0.85699677
-1.03748143]
[ 0.48429942 0.98872018 -0.57314301]
[ 1.52989244 1.16229868 -4.74406147]]
3 8.668 [[ 0.12396193 0.61504567
-0.47498202]
[ 0.22003263 -0.2470119 0.9268558 ]
[ 0.96035379 0.41933775 -3.43156195]]
4 5.77111 [[-0.9524312 1.13037777
0.08607888]
[-3.78651619 2.26245379 2.42393875]
[-3.07170963 3.14037919 -2.12054014]]
5 inf [[ nan nan nan ]
[ nan nan nan ]
[ nan nan nan ]]
6 nan [[ nan nan nan ]
[ nan nan nan ]
[ nan nan nan ]]
...
Prediction: [0 0 0]
Accuracy: 0.0
```

```

X = tf.placeholder("float", [None, 3])
Y = tf.placeholder("float", [None, 3])
W = tf.Variable(tf.random_normal([3, 3]))
b = tf.Variable(tf.random_normal([3]))
hypothesis = tf.nn.softmax(tf.matmul(X, W)+b)
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), reduction_indices=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=1e-10).minimize(cost)

# Correct prediction Test model
prediction = tf.argmax(hypothesis, 1)
is_correct = tf.equal(prediction, tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))

# Launch graph
with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())
    for step in range(201):
        cost_val, W_val, _ = sess.run([cost, W, optimizer],
                                      feed_dict={X: x_data, Y: y_data})
        print(step, cost_val, W_val)

    # predict
    print("Prediction:", sess.run(prediction, feed_dict={X: x_test}))
    # Calculate the accuracy
    print("Accuracy: ", sess.run(accuracy, feed_dict={X: x_test, Y: y_test}))

```

# Small learning rate


```

0 5.73203 [[ 0.80269563  0.67861295 -1.21728313]
[-0.3051686 -0.3032113  1.50825703]
[ 0.75722361 -0.7008909 -2.10820389]]
1 5.73203 [[ 0.80269563  0.67861295 -1.21728313]
[-0.3051686 -0.3032113  1.50825703]
[ 0.75722361 -0.7008909 -2.10820389]]
...
198 5.73203 [[ 0.80269563  0.67861295 -1.21728313]
[-0.3051686 -0.3032113  1.50825703]
[ 0.75722361 -0.7008909 -2.10820389]]
199 5.73203 [[ 0.80269563  0.67861295 -1.21728313]
[-0.3051686 -0.3032113  1.50825703]
[ 0.75722361 -0.7008909 -2.10820389]]
200 5.73203 [[ 0.80269563  0.67861295 -1.21728313]
[-0.3051686 -0.3032113  1.50825703]
[ 0.75722361 -0.7008909 -2.10820389]]
Prediction: [0 0 0]
Accuracy: 0.0

```

# Non-normalized inputs

```
xy = np.array([[828.659973, 833.450012, 908100, 828.349976, 831.659973],  
[823.02002, 828.070007, 1828100, 821.655029, 828.070007],  
[819.929993, 824.400024, 1438100, 818.97998, 824.159973],  
[816, 820.958984, 1008100, 815.48999, 819.23999],  
[819.359985, 823, 1188100, 818.469971, 818.97998],  
[819, 823, 1198100, 816, 820.450012],  
[811.700012, 815.25, 1098100, 809.780029, 813.669983],  
[809.51001, 816.659973, 1398100, 804.539978, 809.559998]])
```



# Non-normalized inputs

```
xy=...
x_data = xy[:, 0:-1]
y_data = xy[:, [-1]]
# placeholders for a tensor that will be always fed.
X = tf.placeholder(tf.float32, shape=[None, 4])
Y = tf.placeholder(tf.float32, shape=[None, 1])
W = tf.Variable(tf.random_normal([4, 1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')

hypothesis = tf.matmul(X, W) + b
cost = tf.reduce_mean(tf.square(hypothesis - Y))
# Minimize
optimizer = tf.train.GradientDescentOptimizer(learning_rate=1e-5)
train = optimizer.minimize(cost)

sess = tf.Session()
sess.run(tf.global_variables_initializer())
for step in range(2001):
    cost_val, hy_val, _ = sess.run(
        [cost, hypothesis, train], feed_dict={X: x_data, Y: y_data})
    print(step, "Cost: ", cost_val, "\nPrediction:\n", hy_val)
```

5 Cost: inf

Prediction:

[[ inf]

[ inf]

[ inf]

...

6 Cost: nan

Prediction:

[[ nan]

[ nan]

[ nan]


...

# Normalized inputs (min-max scale)

```
xy = np.array([[828.659973, 833.450012, 908100, 828.349976, 831.659973],  
[823.02002, 828.070007, 1828100, 821.655029, 828.070007],  
[819.929993, 824.400024, 1438100, 818.97998, 824.159973],  
[816, 820.958984, 1008100, 815.48999, 819.23999],  
[819.359985, 823, 1188100, 818.469971, 818.97998],  
[819, 823, 1198100, 816, 820.450012],  
[811.700012, 815.25, 1098100, 809.780029, 813.669983],  
[809.51001, 816.659973, 1398100, 804.539978, 809.559998]])
```

```
[[ 0.9999999  0.9999999  0.          1.          1.        ]  
[ 0.70548491 0.70439552 1.          0.71881782 0.83755791]  
[ 0.54412549 0.50274824 0.57608696 0.606468  0.6606331 ]  
[ 0.33890353 0.31368023 0.10869565 0.45989134 0.43800918]  
[ 0.51436   0.42582389 0.30434783 0.58504805 0.42624401]  
[ 0.49556179 0.42582389 0.31521739 0.48131134 0.49276137]  
[ 0.11436064 0.          0.20652174 0.22007776 0.18597238]  
[ 0.          0.07747099 0.5326087  0.          0.        ]]
```

```
xy = MinMaxScaler(xy)  
print(xy)
```



# Normalized inputs

```
xy=...
x_data = xy[:, 0:-1]
y_data = xy[:, [-1]]
# placeholders for a tensor that will be always fed.
X = tf.placeholder(tf.float32, shape=[None, 4])
Y = tf.placeholder(tf.float32, shape=[None, 1])
W = tf.Variable(tf.random_normal([4, 1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')

hypothesis = tf.matmul(X, W) + b
cost = tf.reduce_mean(tf.square(hypothesis - Y))
# Minimize
optimizer = tf.train.GradientDescentOptimizer(learning_rate=1e-5)
train = optimizer.minimize(cost)

sess = tf.Session()
sess.run(tf.global_variables_initializer())
for step in range(2001):
    cost_val, hy_val, _ = sess.run(
        [cost, hypothesis, train], feed_dict={X: x_data, Y: y_data})
    print(step, "Cost: ", cost_val, "\nPrediction:\n", hy_val)
```

Prediction:  
[[ 1.63450289]  
[ 0.06628087]  
[ 0.35014752]  
[ 0.67070574]  
[ 0.61131608]  
[ 0.61466062]  
[ 0.23175186]  
[-0.13716528]]

With TF 1.0!




# Lab 7-2

## MNIST data

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/>

 <b>zeran4</b> 1 commit / 5 ++ / 4 --	#11	 <b>jennykang</b> 19 commits / 940 ++ / 253 --	#2
 <b>GzuPark</b> 14 commits / 41 ++ / 31 --	#3	 <b>kkweon</b> 5 commits / 372 ++ / 296 --	#4
 <b>BlueMelon715</b> 4 commits / 45 ++ / 34 --	#5	 <b>jin-chong</b> 2 commits / 4 ++ / 4 --	#6
 <b>FuZer</b> 2 commits / 37 ++ / 30 --	#7	 <b>cynthia</b> 1 commit / 28 ++ / 28 --	#8
 <b>keon</b> 1 commit / 3 ++ / 3 --	#9	 <b>allieus</b> 1 commit / 55 ++ / 59 --	#10

# MNIST Dataset




[train-images-idx3-ubyte.gz](#): training set images (9912422 bytes)

[train-labels-idx1-ubyte.gz](#): training set labels (28881 bytes)

[t10k-images-idx3-ubyte.gz](#): test set images (1648877 bytes)

[t10k-labels-idx1-ubyte.gz](#): test set labels (4542 bytes)

# 28x28x1 image



```
# MNIST data image of shape 28 * 28 = 784
X = tf.placeholder(tf.float32, [None, 784])
# 0 - 9 digits recognition = 10 classes
Y = tf.placeholder(tf.float32, [None, nb_classes])
```

```
00000000000000000000  
11111111111111111111  
22222222222222222222  
33333333333333333333  
44444444444444444444  
55555555555555555555  
66666666666666666666  
77777777777777777777  
88888888888888888888  
99999999999999999999
```

# MNIST Dataset

```
from tensorflow.examples.tutorials.mnist import input_data  
# Check out https://www.tensorflow.org/get\_started/mnist/beginners for  
# more information about the mnist dataset  
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)  
...  
batch_xs, batch_ys = mnist.train.next_batch(100)  
...  
print("Accuracy: ", accuracy.eval(session=sess,  
    feed_dict={X: mnist.test.images, Y: mnist.test.labels}))
```

# Reading data and set variables

```
from tensorflow.examples.tutorials.mnist import input_data
# Check out https://www.tensorflow.org/get_started/mnist/beginners for
# more information about the mnist dataset
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)

nb_classes = 10

# MNIST data image of shape 28 * 28 = 784
X = tf.placeholder(tf.float32, [None, 784])
# 0 - 9 digits recognition = 10 classes
Y = tf.placeholder(tf.float32, [None, nb_classes])

W = tf.Variable(tf.random_normal([784, nb_classes]))
b = tf.Variable(tf.random_normal([nb_classes]))
```

# Softmax!

```
# Hypothesis (using softmax)
hypothesis = tf.nn.softmax(tf.matmul(X, W) + b)

cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)

# Test model
is_correct = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))
# Calculate accuracy
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

# Training epoch/batch

```
# parameters
training_epochs = 15
batch_size = 100

with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())
    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0
        total_batch = int(mnist.train.num_examples / batch_size)

        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
            c, _ = sess.run([cost, optimizer], feed_dict={X: batch_xs, Y: batch_ys})
            avg_cost += c / total_batch

        print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))
```

# Training epoch/batch

In the neural network terminology:

- one **epoch** = one forward pass and one backward pass of *all* the training examples
- **batch size** = the number of training examples in one forward/backward pass. The higher the batch size, the more memory space you'll need.
- number of **iterations** = number of passes, each pass using [batch size] number of examples. To be clear, one pass = one forward pass + one backward pass (we do not count the forward pass and backward pass as two different passes).

Example: if you have *1000 training examples*, and your *batch size is 500*, then it will take 2 iterations to complete 1 epoch.

# Training epoch/batch

```
# parameters
training_epochs = 15
batch_size = 100

with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())
    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0
        total_batch = int(mnist.train.num_examples / batch_size)

        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
            c, _ = sess.run([cost, optimizer], feed_dict={X: batch_xs, Y: batch_ys})
            avg_cost += c / total_batch

        print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))
```

# Report results on test dataset

```
# Test the model using test sets
print("Accuracy: ", accuracy.eval(session=sess,
    feed_dict={X: mnist.test.images, Y: mnist.test.labels}))
```

```

hypothesis = tf.nn.softmax(tf.matmul(X, W) + b)
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)

is_correct = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))

# parameters
training_epochs = 15
batch_size = 100

with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())
    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0
        total_batch = int(mnist.train.num_examples / batch_size)

        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
            c, _ = sess.run([cost, optimizer],
                           feed_dict={X: batch_xs, Y: batch_ys})
            avg_cost += c / total_batch

        print('Epoch:', '%04d' % (epoch + 1),
              'cost =', '{:.9f}'.format(avg_cost))


```

```

Epoch: 0001 cost = 2.868104637
Epoch: 0002 cost = 1.134684615
Epoch: 0003 cost = 0.908220728
Epoch: 0004 cost = 0.794199896
Epoch: 0005 cost = 0.721815854
Epoch: 0006 cost = 0.670184430
Epoch: 0007 cost = 0.630576546
Epoch: 0008 cost = 0.598888191
Epoch: 0009 cost = 0.573027079
Epoch: 0010 cost = 0.550497213
Epoch: 0011 cost = 0.532001859
Epoch: 0012 cost = 0.515517795
Epoch: 0013 cost = 0.501175288
Epoch: 0014 cost = 0.488425370
Epoch: 0015 cost = 0.476968593
Learning finished
Accuracy: 0.888

```

# Sample image show and prediction



The screenshot shows a Jupyter Notebook interface with two cells. The left cell displays a 28x28 pixel grayscale image of a handwritten digit '6' on a white background. The right cell contains the following Python code:

```
import matplotlib.pyplot as plt
import random

# Get one and predict
r = random.randint(0, mnist.test.num_examples - 1)
print("Label:", sess.run(tf.argmax(mnist.test.labels[r:r+1], 1)))
print("Prediction:", sess.run(tf.argmax(hypothesis, 1),
                             feed_dict={X: mnist.test.images[r:r + 1]}))

plt.imshow(mnist.test.images[r:r + 1].
           reshape(28, 28), cmap='Greys', interpolation='nearest')
plt.show()
```

The notebook's toolbar and status bar are visible at the bottom, showing the current file is 'lab-07-4-mnist\_introduction.py'.

# Lab 8

## Tensor Manipulation

Sung Kim <hunkim+ml@gmail.com>